

72 hr

We are building:

OmniQuant v2 — Arbitrage Detection & Market Microstructure Simulation Platform

◆ GLOBAL RULES (Before You Start)

1. Stability > Features
 2. Core engine must work before UI
 3. Every module must produce measurable output
 4. All results must be labeled “Simulated / Theoretical”
 5. No real trade execution
-

🧭 PHASE 1 — FOUNDATION & ARCHITECTURE SETUP

Objective:

Create the structural backbone of the system.

Deliverables:

- Repository structure
 - Build system working
 - C++ core compiling
 - Python bridge connected
 - Basic API returning mock data
-

Step 1 — Create Repository Structure

```
OmniQuant/ | └── core/ (C++ Engine) └── bindings/ (pybind11) └── simulation/  
|   └── risk/ └── optimizer/ └── analytics/ └── api/ └── frontend/ └── tests/ └──  
README.md
```

Step 2 — Setup C++ Build

Use:

- CMake
- C++17

Ensure:

- Builds clean
 - Outputs shared library
-

Step 3 — Setup Python Bridge

Use:

- pybind11

Expose minimal function:

```
detect_cycles(market_data)
```

Return dummy cycle first.

Test:

```
import omniquant_core print(omniquant_core.detect_cycles(...))
```

If this works → foundation is stable.



PHASE 2 — CORE ARBITRAGE ENGINE (C++)

Objective:

Make arbitrage detection mathematically correct.

Step 1 — Graph Model

- Directed weighted graph
- Weight = $-\log(\text{rate} \times (1 - \text{fee}))$

Store:

- Nodes
 - Edges
 - Metadata
-

Step 2 — Negative Cycle Detection

Implement:

- Bellman-Ford
- Cycle reconstruction

Output:

- Path
- Raw theoretical return

Test with synthetic data:

- Force known arbitrage cycle
-

Step 3 — Edge Pruning

Before running Bellman-Ford:

- Remove low-liquidity edges
- Remove extreme-fee edges

Improves performance and realism.

Step 4 — Performance Metrics

Track:

- Detection runtime (ms)
- Graph size
- Edge count

Expose to Python.



PHASE 3 — MARKET MICROSTRUCTURE SIMULATION (PYTHON)

Objective:

Turn raw arbitrage into realistic execution model.

Step 1 — Order Book Model

For each pair:

- Simulated top 5 bids/asks
 - Associated depth
-

Step 2 — Market Impact Model

Implement nonlinear impact:

$$\text{impact} = k \cdot (\text{volume}/\text{liquidity})^\alpha$$
$$\text{impact} = k \cdot (\text{volume}/\text{liquidity})^{\alpha}$$
$$\text{impact} = k \cdot (\text{volume}/\text{liquidity})^\alpha$$

Apply to execution price.

Step 3 — Slippage Model

Compute:

- Effective rate after depth consumption
 - Multi-hop compounding impact
-

Step 4 — Latency Sensitivity

Simulate:

- 0ms
- 50ms
- 100ms delay

Adjust price drift randomly.

Compute:

- Arbitrage half-life
-

Step 5 — Monte Carlo Engine

For each opportunity:

- Run 500–1000 simulations
- Randomize:
 - Latency
 - Slippage
 - Minor volatility noise

Return:

- Mean return
- Std deviation
- 5% worst case
- Probability negative

This makes system institutional-grade.

🛡 PHASE 4 — RISK ENGINE

Objective:

Quantify uncertainty.

Step 1 — Liquidity Risk Score

Based on:

- Volume consumed / depth
-

Step 2 — Complexity Risk

Longer cycle → higher risk.

Step 3 — Volatility Exposure

Estimate short-window volatility.

Step 4 — Composite Risk Score

Weighted aggregation:

$$\text{Risk} = w_1L + w_2C + w_3VRisk = w_1L + w_2C + w_3VRisk=w1L+w2C+w3V$$

Normalize 0–100.

Step 5 — Confidence Metric

$$\begin{aligned}\text{Confidence} &= 1 - \text{Probability}(\text{negative_return}) \\ \text{Confidence} &= 1 - \\ \text{Probability}(\text{negative_return}) &\text{Confidence} = 1 - \text{Probability}(\text{negative_return})\end{aligned}$$



PHASE 5 — ANALYTICS & PERSISTENCE

Objective:

Look like a research tool.

Step 1 — Opportunity Tracking

Track:

- First detected timestamp
 - Duration alive
 - Peak return
 - Decay pattern
-

Step 2 — Persistence Metric

Measure:

- Frequency across time window
- Stability of return

Compute:

- Mean
 - Variance
 - Sharpe-like ratio
-

Step 3 — Stress Testing

Simulate:

- +1% uniform price shock
- -1% liquidity shock
- Fee increase

Check survival rate.

PHASE 6 — CAPITAL OPTIMIZER

Objective:

Institutional-level decision logic.

Step 1 — Input:

- Capital budget
 - List of opportunities
 - Risk scores
 - Liquidity constraints
-

Step 2 — Optimization

Maximize:

$$\sum x_i r_i$$

Subject to:

- $\sum x_i \leq \text{capital}$
- Risk budget limit
- Liquidity cap

Start with greedy allocation.

Advanced: linear programming.

PHASE 7 — API LAYER (FASTAPI)

Endpoints:

/scan

Triggers detection pipeline.

/opportunities

Returns ranked list.

/metrics

Returns system telemetry.

/stress

Returns robustness results.

PHASE 8 — FRONTEND DASHBOARD

Minimal but elite.

Sections:

1. Ranked Opportunities Table
2. Risk Breakdown Panel
3. Monte Carlo Distribution Chart
4. Latency Decay Curve
5. System Telemetry

Dark theme. Clean typography.

No clutter.

PHASE 9 — AI EXPLANATION LAYER

Convert structured output into:

- Summary
- Expected return
- Worst-case scenario
- Risk level
- Robustness indicator
- Half-life

Always add disclaimer footer.

PHASE 10 — DOCUMENTATION

README must include:

1. Mathematical modeling explanation
2. Graph-theoretic arbitrage proof
3. Simulation assumptions
4. Risk methodology
5. Known limitations

6. Future real-trade integration plan

This is where judges decide seriousness.

FINAL SYSTEM CHECKLIST

Before demo:

- Arbitrage detection works on synthetic + real API data
 - Monte Carlo produces distribution plot
 - Risk score varies meaningfully
 - Stress test changes results
 - Latency half-life visible
 - Capital allocator outputs suggestion
 - Dashboard responsive
 - Disclaimers visible everywhere
-

WHAT YOU END UP WITH

Not:

“Crypto money machine”

But:

A deterministic arbitrage detection and execution-simulation research engine with statistical validation and microstructure modeling.

That is how a strong MIT/IIT systems + quant student would build it.