

Implementation v2

◆ OmniQuant v2 — Quantitative Market Inefficiency Research Platform

This is now MIT-grade.

CORE PHILOSOPHY SHIFT

Before:

| Detect arbitrage cycles.

Now:

| Model, measure, stress-test, and rank structural inefficiencies under uncertainty.

That is a different level.

NEW FEATURES (JANE STREET MODE)

We keep your previous system.

Now we add layers.

1 Order Book Microstructure Modeling

Your v1 assumes a single price.

Real markets have:

- Bid
- Ask
- Depth
- Spread
- Impact cost

Add:

- ◆ **Level-2 Order Book Simulation**

For each exchange:

- Store top N bids/asks
- Simulate market impact

Profit calculation becomes:

$$P = \prod_i (\text{executed_price}_i) - P_0 = \prod_i (\text{executed_price}_i) - P_0 = \prod_i (\text{executed_price}_i) - P_0$$

Where `executed_price` depends on depth consumed.

2 Market Impact Function

Implement nonlinear slippage:

$$\text{impact} = k \cdot (\text{volume} / \text{liquidity})^\alpha$$
$$\text{impact} = k \cdot \left(\frac{\text{volume}}{\text{liquidity}}\right)^\alpha$$

Where:

- k = sensitivity constant
- $\alpha > 1$ (convex impact)

This shows institutional understanding.

3 Latency Sensitivity Analysis

Instead of one return value, compute:

- Return at 0ms delay
- Return at 50ms
- Return at 200ms

Plot decay curve:

If opportunity disappears in 30ms → low reliability.

Add:

◆ Arbitrage Half-Life Metric

Time until expected return = 0.

That is elite.

4 Statistical Persistence Testing

Don't just detect once.

Run detection over time window:

- Track frequency of cycle
- Track variance of return
- Compute Sharpe-like metric:

$$\text{Sharpe} = \frac{\text{E}[R] - R_f}{\sigma_R}$$

Now judges see quant rigor.

5 Capital Allocation Optimizer

Given multiple cycles:

Maximize:

$$\max \sum_i x_i r_i$$

Subject to:

- Liquidity constraints
- Capital limit
- Risk budget

Solve via:

- Linear programming (Python)
- Or simple greedy approximation

This turns OmniQuant into portfolio allocator.

6 Risk Engine (Real One)

Add risk components:

- ◆ **Liquidity Risk**
- ◆ **Path Complexity Risk**
- ◆ **Volatility Exposure**
- ◆ **Execution Uncertainty**
- ◆ **Spread Risk**

Aggregate via weighted model:

$$\text{RiskScore} = w_1L + w_2V + w_3C + w_4S \quad \text{RiskScore} = w_1L + w_2V + w_3C + w_4S \quad \text{RiskScore} = w_1L + w_2V + w_3C + w_4S$$

Return:

- Risk percentile
- Confidence band

7 Stress Testing Module

Simulate shocks:

- ±1% price move
- Liquidity drop 30%
- Fee increase

Check:

Does arbitrage survive?

Output:

| Stress Robustness: 42%

This is hedge-fund thinking.

8 Opportunity Ranking Engine

Rank opportunities by:

$$\text{Score} = \frac{\text{ExpectedReturn} \cdot \text{Confidence}}{\text{RiskScore}}$$

Score = RiskScore / ExpectedReturn * Confidence

Display leaderboard.

9 Regime Detection

Classify market regime:

- High volatility
- Low liquidity
- Stable

Use simple volatility clustering.

Adjust risk model accordingly.

That is advanced thinking.

10 Performance Telemetry Dashboard

Add:

- Detection time (ms)
- Graph size
- CPU usage
- Cycle search complexity

Engineers respect instrumentation.

11 Monte Carlo Execution Simulator

For each opportunity:

Simulate 1,000 runs with random:

- Slippage
- Latency
- Volatility noise

Return distribution:

- Mean return
- Worst 5% outcome
- Probability of negative outcome

This makes it academic-level.

1|2 Edge Attribution Analysis

Break down profit source:

- Fee inefficiency %
- Spread mispricing %
- Multi-hop distortion %

This shows analytical depth.

1|3 Opportunity Lifecycle Tracking

Track:

- When detected
- Duration alive
- Peak return
- Decay pattern

Display:

Opportunity lifespan chart.

1|4 Dynamic Graph Pruning

Instead of full graph always:

- Prune low-liquidity edges
- Remove low-confidence edges
- Adaptive edge weighting

Reduces complexity.

Very MIT.

1 | 5 Modular Engine Architecture

Refactor:

```
/core graph_engine.cpp cycle_detector.cpp edge_pruner.cpp /simulation  
slippage_model.py impact_model.py monte_carlo.py /risk risk_engine.py  
stress_test.py /optimizer capital_allocator.py /analytics  
persistence_tracker.py regime_detector.py /api main.py /frontend dashboard
```

This looks like research software.

FINAL OMNIQUANT V2 IMPLEMENTATION BLUEPRINT

Follow in this order.

PHASE 1 — Core Engine (C++)

- ✓ Graph construction
 - ✓ Negative cycle detection
 - ✓ Path extraction
 - ✓ Edge pruning logic
 - ✓ Performance metrics
-

PHASE 2 — Advanced Simulation Layer (Python)

- ✓ Level-2 order book model
- ✓ Nonlinear impact function
- ✓ Latency sensitivity curve

- ✓ Monte Carlo simulator
 - ✓ Stress test module
-

PHASE 3 — Risk & Analytics Layer

- ✓ Risk scoring engine
 - ✓ Confidence interval computation
 - ✓ Persistence tracker
 - ✓ Sharpe-like metric
 - ✓ Regime classification
-

PHASE 4 — Optimization Layer

- ✓ Capital allocator
 - ✓ Opportunity ranking
 - ✓ Risk-budget constraint
-

PHASE 5 — AI Explanation Layer

Now explanation includes:

- Expected return
- Worst-case simulated return
- Risk percentile
- Robustness score
- Half-life metric

Example:

OmniQuant detected a multi-hop pricing inefficiency.
Expected return: 0.21%
Monte Carlo worst-case: -0.05%
Latency half-life: 42ms
Risk percentile: 61%
Stress robustness: Moderate

That is MIT-grade.

PHASE 6 — Dashboard

Sections:

1. Live Opportunities
2. Ranked Leaderboard
3. Risk Breakdown
4. Monte Carlo Distribution Plot
5. Latency Decay Curve
6. System Telemetry

Clean. Minimal. Quant aesthetic.

PHASE 7 — Mandatory Disclaimers

Add persistent banner:

OmniQuant is a research-grade arbitrage detection and simulation platform.
No trades are executed.
Results are theoretical and subject to model assumptions.
Not financial advice.

Never remove this.

What This Becomes

Not a hackathon toy.

But:

A Market Microstructure Research Engine With Explainable Intelligence.

That is something a strong IIT CSE or MIT student would proudly demo.